

# PROCESS MANAGEMENT

---

# PREVIOUS VIDEO

- Introduction to operating system
- Kernel vs operating system
- Functions of operating system
- Examples of operating system and kernel
- System calls
- Interrupts
- Video Link + telegram link + form link in the description

# AGENDA

- Program VS process
- Process control block
- Process states
- Process creation through fork()
- Exec(), wait() and exit() system calls
- Zombie and orphan process

# INTERVIEW QUESTIONS

- Difference between a program and a process
- Which data structure is used to represent a process?
- List some important fields in Process Control Block
- What is INIT process in Linux?
- What are the different states of a process? Explain with process state diagram.
- How will you create a new process in Linux?
- How does a parent process differ from a child process?
- Number of processes created with N fork() calls?

# INTERVIEW QUESTIONS CONTINUED

- What is the use of `exec()`, `wait()` and `exit()` system calls?
- What happens if parent doesn't wait for its child process?
- Give some real world examples where `fork()`, `exec()` and `wait()` system calls are used?
- What is zombie process? What are the side effects of zombie processes?
- What is an orphan process? How it differs from zombie process?
- Give examples of zombie and orphan process.

# PROGRAM VS PROCESS

- Program – An executable file that is stored on the hard disk or a secondary memory
- When you execute this file either by double clicking or from command line, it becomes a process
- Program – Passive entity
- Process – Active entity
- Process memory = stack + data + heap + text
- Program is just the text section of memory
- Process is text section + a lot of other attributes

# What are other attributes of a process?

- Program counter(PC) – Address of next instruction to be executed on CPU
- Process states – new, ready, running, waiting etc.
- CPU registers – value of CPU registers are stored when context switch happens
- List of open files that the process is using for reading and writing
- Process priority – Used in process scheduling
- Process ID – unique identifier for each process
- Parent process ID – Process ID of parent process
- Memory information – stack, heap, data, text

# PROCESS CONTROL BLOCK

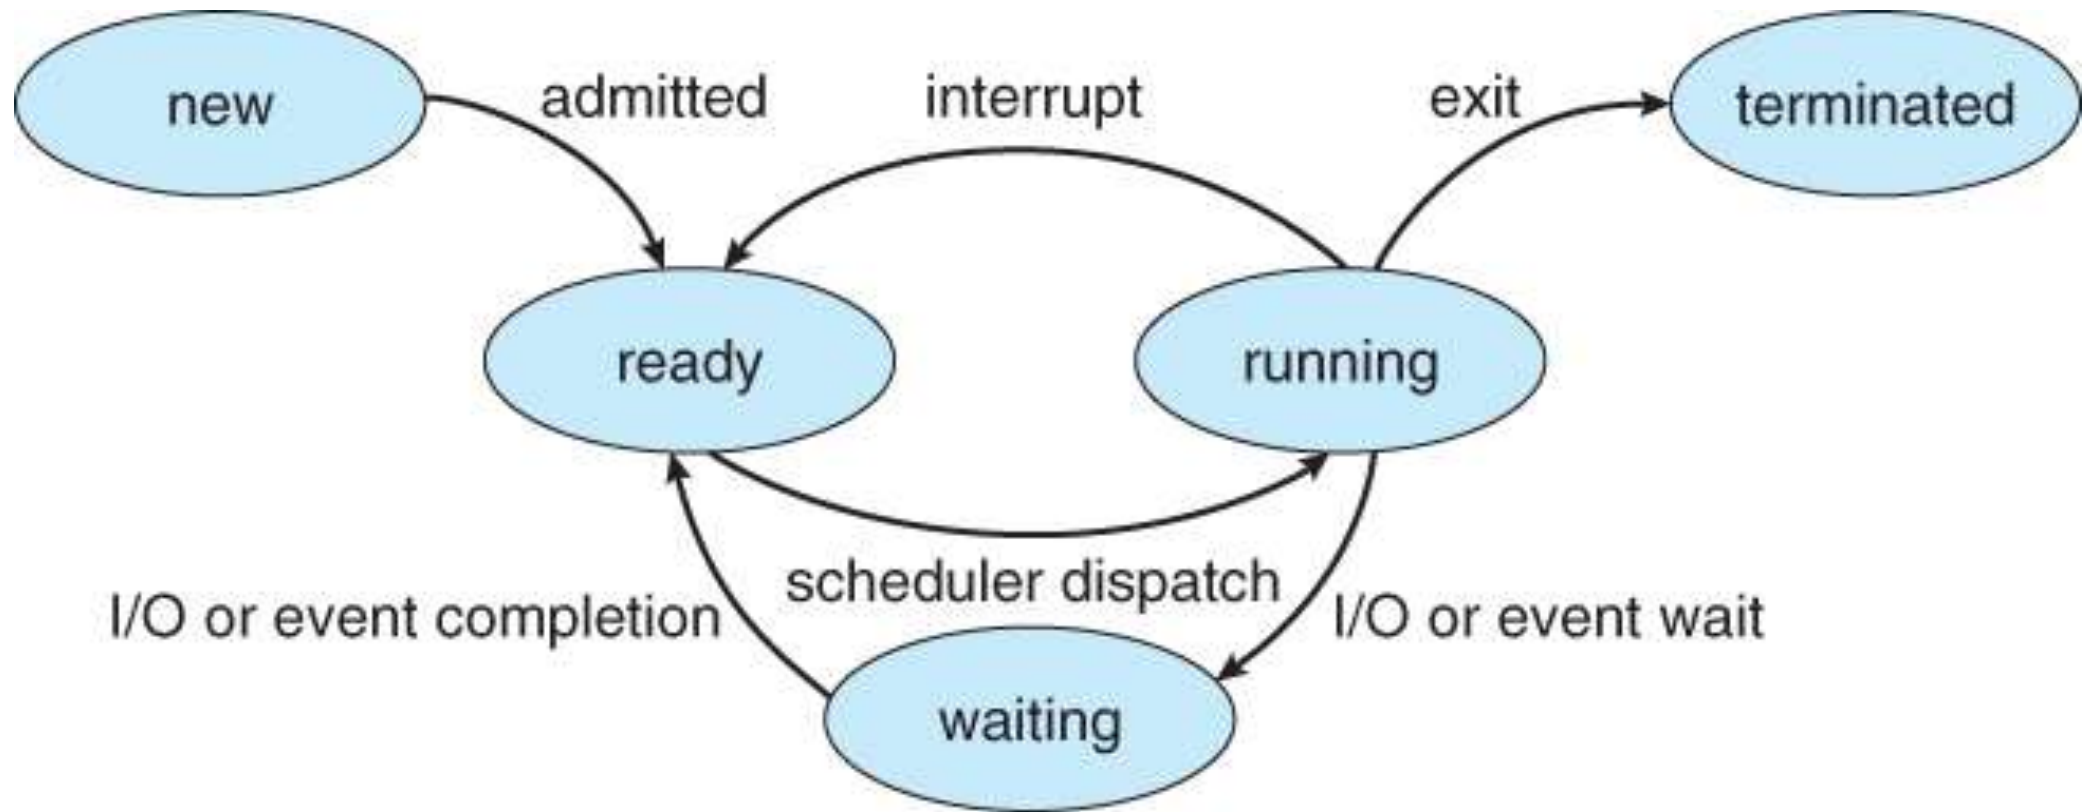
- These all attributes or information about a process are stored in a data structure called Process Control Block(PCB).
- Struct task\_struct {
  - pid\_t pid;
  - long state;
  - int priority;
  - struct task\_struct \*parent;
  - struct files\_struct \*files;
  - struct mm\_struct \*mm;
  - long program\_counter;
  - .
  - .}
- Video: <https://youtu.be/x2JIEgAVkTI>



# PROCESS STATES

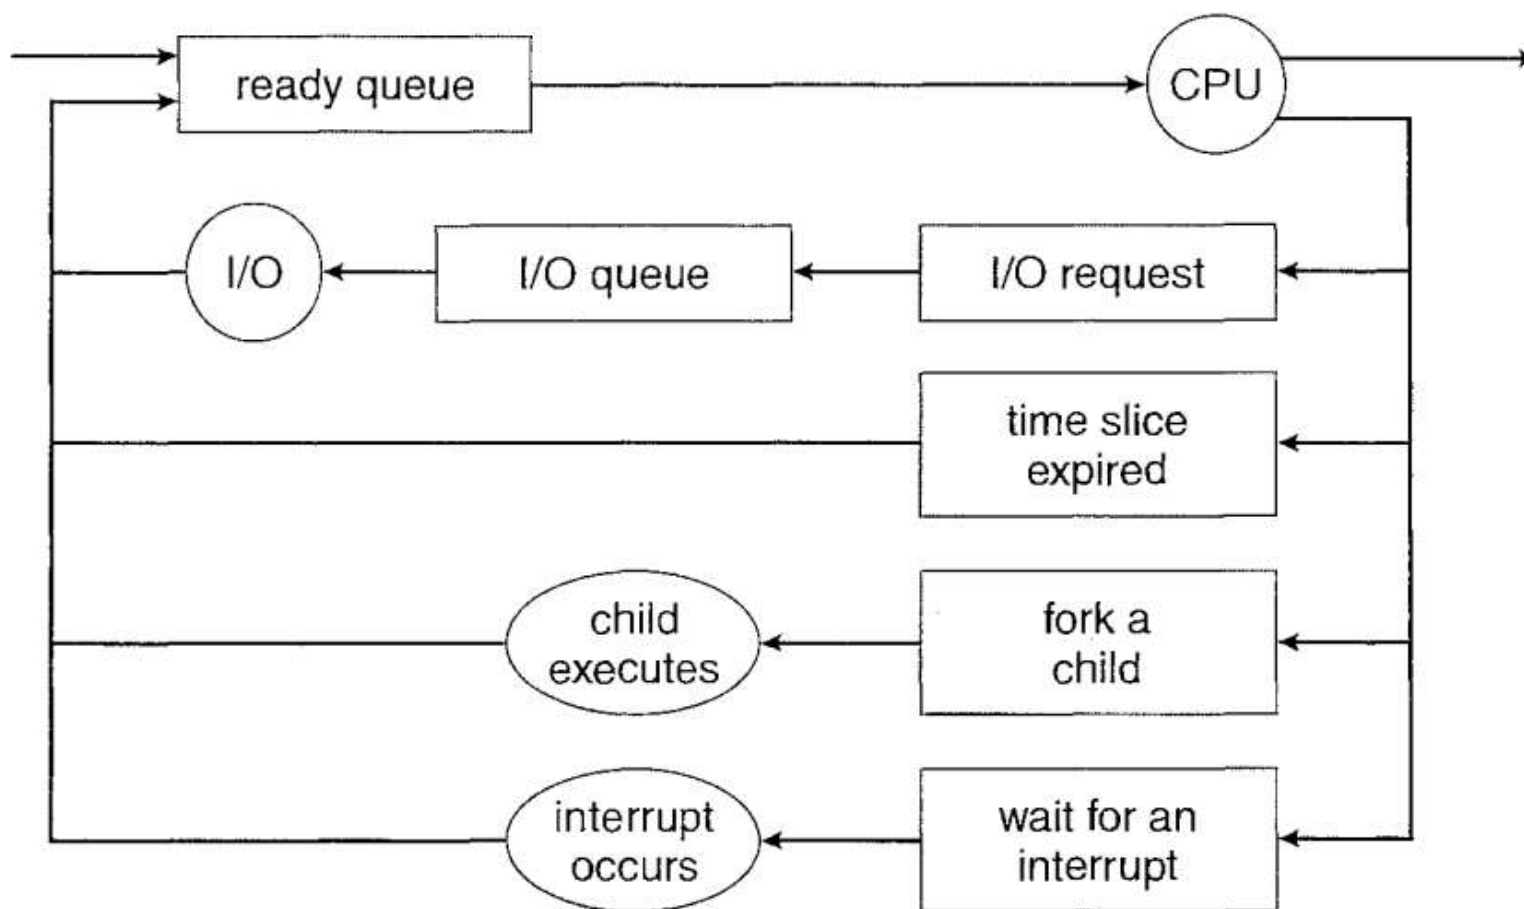
- Batch operating system in old times
- Jobs were submitted in batches
- Jobs were first stored in job queue on disk
- Not all job can fit in main memory
- Selected jobs were loaded into main memory by long term scheduler
- These were kept in ready queue and scheduled by CPU scheduler or short term scheduler
- Video: <https://youtu.be/mhKr6Dnf2fA>

# PROCESS STATE DIAGRAM



# WHAT EACH STATE MEANS?

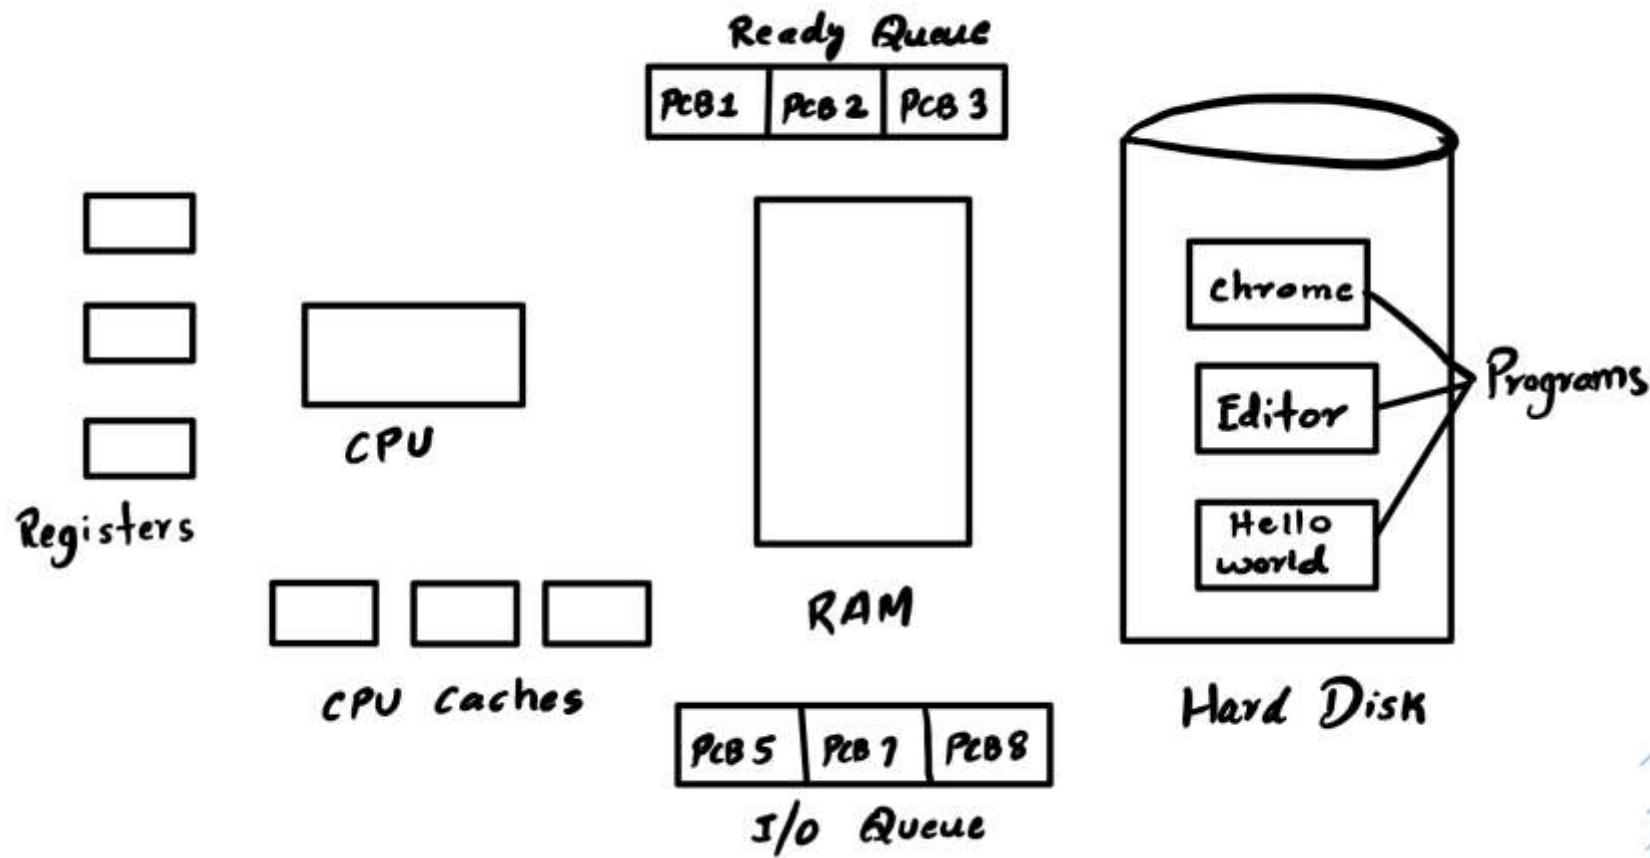
- NEW
  - BATCH OS – Job arrived in the job queue
  - Time sharing OS – process is being created
- READY – Process waiting in ready queue for being assigned to a CPU/processor
- RUNNING – Program is being executed on the CPU
- WAITING – Process is waiting for some event to occur (I/O)
- TERMINATED – Process has finished its execution.



**Figure 3.7** Queueing-diagram representation of process scheduling.

# WHEN DOES TRANSITION FROM RUNNING TO READY HAPPENS?

- A hardware interrupt happens on the CPU
- A higher priority process arrives in the ready queue (preemptive scheduling)
- Time slice of the process expires (Round robin scheduling)



- Program – helloworld.c
- Compile – gcc helloworld.c
- Executable – a.out
- Execute the program
- a.out is bought to RAM from disk
- A PCB is created (PC -> address of first instruction )
- Enters ready queue
- CPU scheduler schedules it on CPU
- Running state
- CPU starts executing
- I/O – enters waiting state
- Time slice expires – ready state
- Terminated

# INIT PROCESS

- First process that is started during booting of the system
- PID – 1
- This process keeps on running forever until you shut down.
- All other processes are either direct or indirect children of INIT process.
- Pstree –p command to view the process tree on linux

# PROCESS CREATION USING FORK

- CreateProcess() in Windows
- Fork() creates a new child process that is a copy of the parent process
- In what ways the child process differs from parent?
  - PID
  - Parent PID
  - File locks
  - Process utilization info
- In what ways they are similar?
  - Program counter
  - Memory ( unless fork() is followed by exec() )
  - CPU registers
  - priority



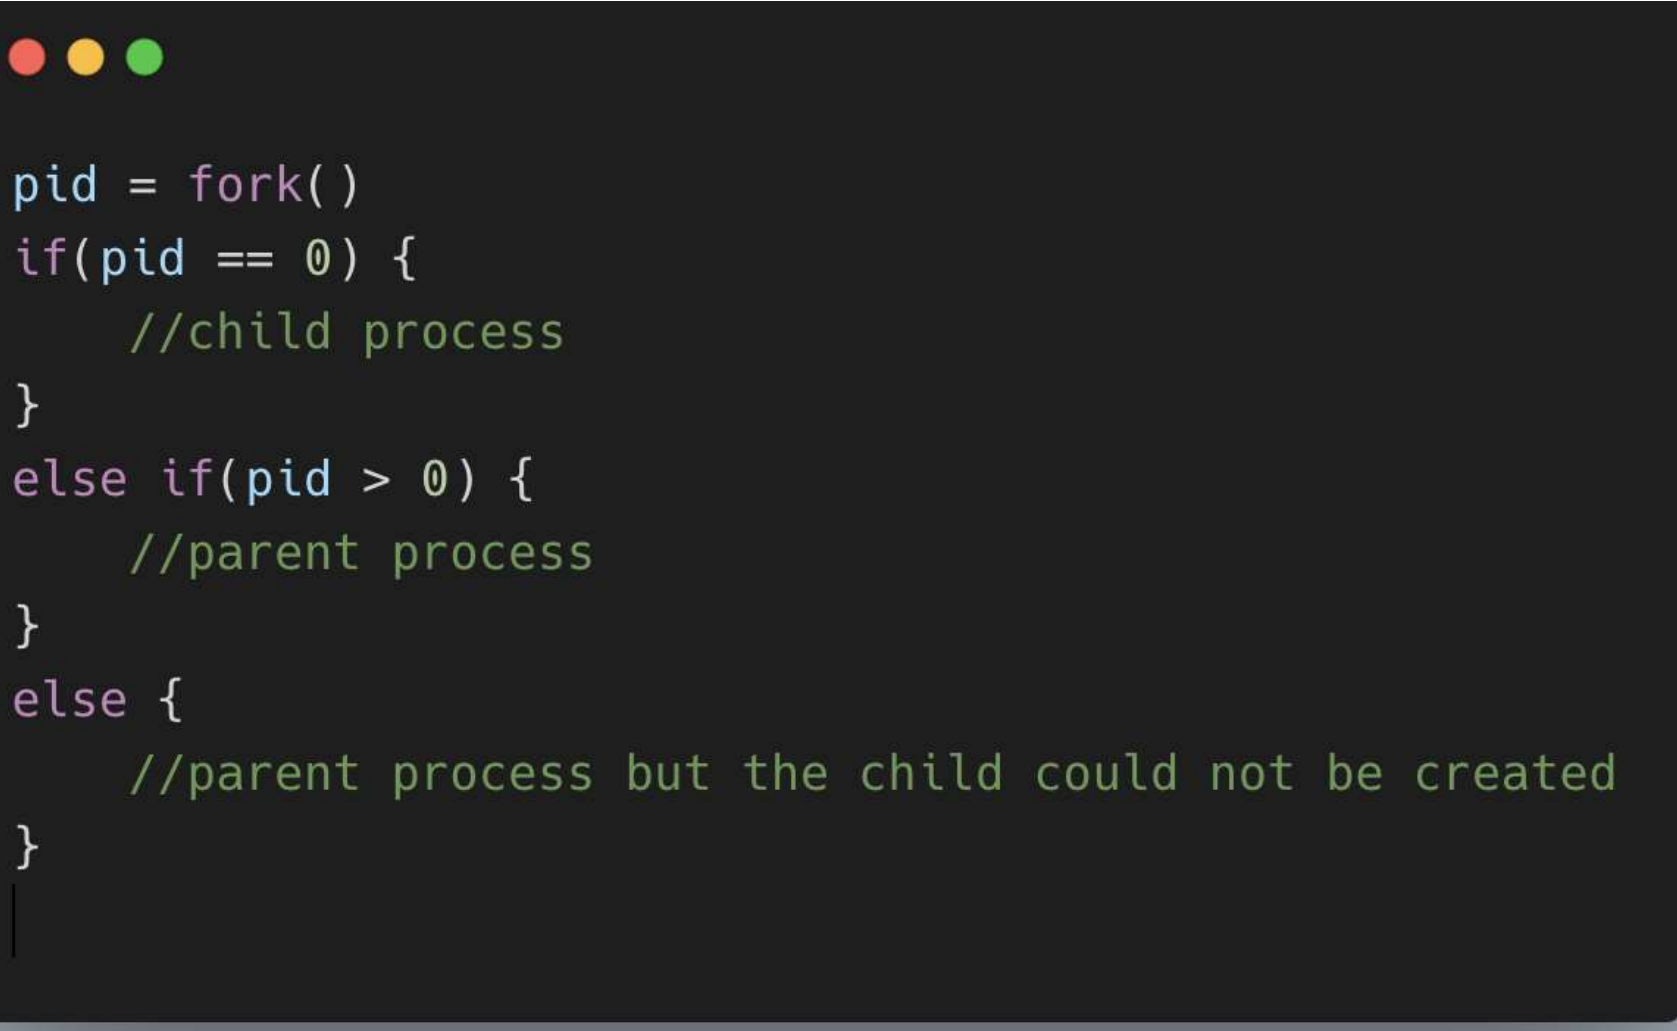
# WHAT HAPPENS AFTER FORK() IS CALLED?

- Execution – Both process starts executing from the next line after fork()
- Memory – stack, heap, text and data segment are duplicated for the child process.
- Copy on write is used in modern implementation of fork.
- Copy the memory only when a write is being made to a memory page.
- In most cases, fork() is followed by exec().
- Copy on write saves expensive recopying of memory

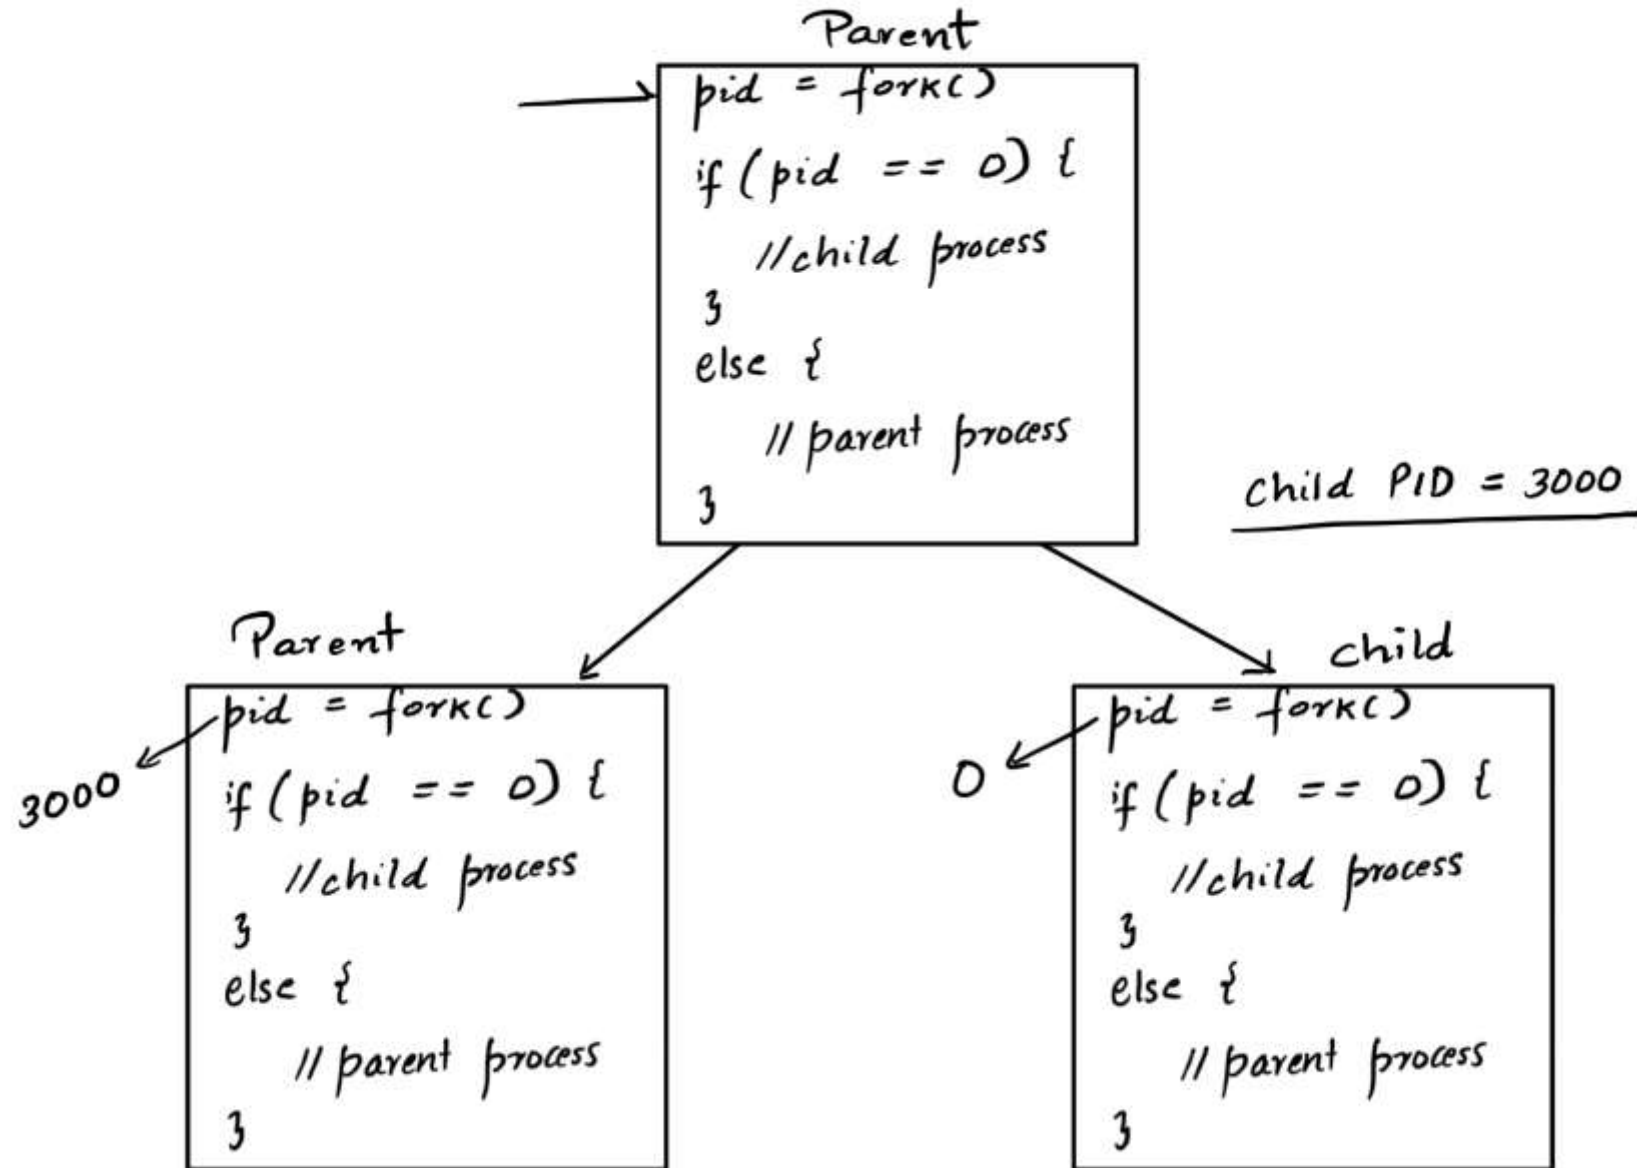
# FORK CALL RETURN VALUES

- How do you differentiate between the the parent process and the child process?
- Fork() returns PID of the child process to parent
- Fork() returns 0 to the child process
- Fork() returns -1 if there is any error
- Videos
  - <https://youtu.be/FXAvkNY1dGQ>
  - <https://youtu.be/AyZeHBPKdMs>

# FORK CODE EXAMPLE



```
pid = fork()
if(pid == 0) {
    //child process
}
else if(pid > 0) {
    //parent process
}
else {
    //parent process but the child could not be created
}
|
```



# How many child processes?

```
fork();
```

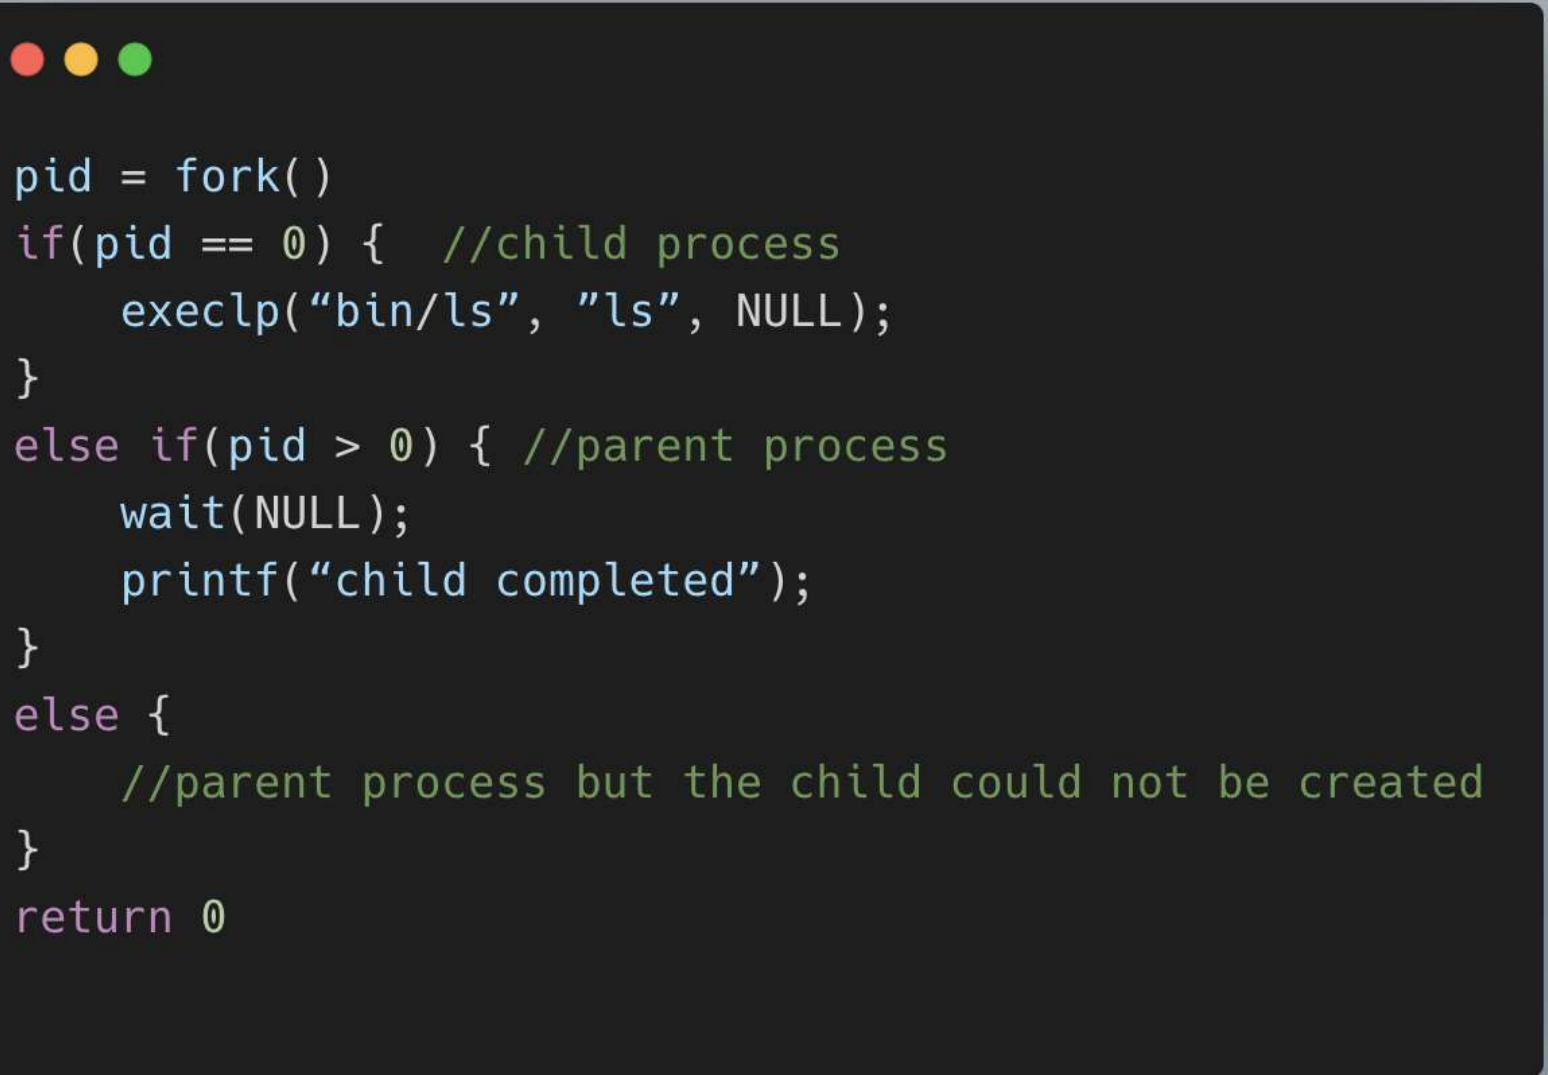
```
fork();
```

Video: <https://youtu.be/iZa2vm7A6mw>

# What if you want the child to run another program?

- Shell program which helps in running different commands on command line
- Thousands of command will be there. Is it a good idea to put all these commands in the shell program?
- What if you want to run a program from shell that is not present in shell program? ( Running your C program )
- The shell program typically uses `fork()` to create a new child process and asks the child process to run the command.

# EXEC(), WAIT() AND EXIT() SYSTEM CALLS

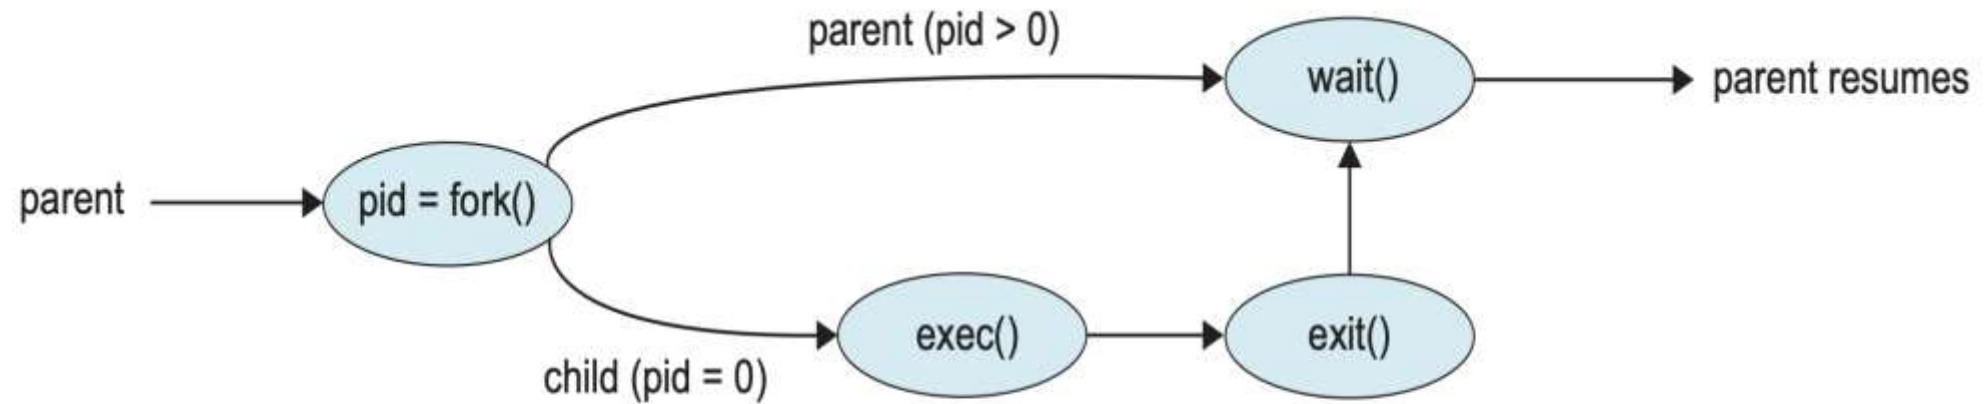


```
pid = fork()
if(pid == 0) { //child process
    execlp("bin/ls", "ls", NULL);
}
else if(pid > 0) { //parent process
    wait(NULL);
    printf("child completed");
}
else {
    //parent process but the child could not be created
}
return 0
```

# What each system call does?

- `exec()` – loads a new program into the main memory and executes it.
- A process exits or terminates using `exit()` system call. The resources like memory/cpu of the process are set free.
- return status statement from main function = `exit(status)`
- PCB of the process is still present in memory
- `wait()` – the parent waits for the completion of the child process. Parent enters the wait state.
- The parent may be interested in knowing about the status of the child.
- After getting the status, the parent asks the kernel to clean up the PCB of child process.





**Figure 3.10** Process creation using the `fork()` system call.

# What happens if the parent doesn't wait for the child process?

- The process control block of the child remains in the process table even after it has terminated using `exit()`.
- These processes are called zombie process - process has terminated but PCB is still there in process table.
- Once the parent terminates, these zombie process also becomes an orphan process.
- No parent + PCB still in the process table after terminating

# SIDE EFFECTS OF ZOMBIE PROCESS


```
while(1) {  
    fork();  
}
```

- Fork bomb
- Number of PIDs is limited in the operating system. Once all PIDs are taken by the zombie process, no new process can be created.
- Since no new process can be created, one cannot even run a command on the command line.
- The only option to recover from a fork() bomb is to reboot the system.

# Orphan Process


- A process which doesn't have a parent process is called an orphan process.
- This may happen if the parent process terminates before child completes.
- What happens to the orphan process?
- Orphan process gets reparented and the new parent is the INIT process in most cases.
- The new parent then waits for the child to complete and then asks the kernel to clean the PCB.

# ZOMBIE PROCESS PROGRAM



```
#include <stdlib.h>
#include <unistd.h>
int main()
{
    int pid = fork();
    if (pid > 0) {
        // Parent process code
        sleep(60);
    }
    if(pid == 0) {
        //child process code
        exit(0);
    }
}
```

# ORPHAN PROCESS PROGRAM

```
  
#include <unistd.h>  
#include <stdlib.h>  
int main()  
{  
    int pid = fork();  
    if (pid > 0) {  
        // parent process code  
        exit(0); // Parent terminates before the child process  
    }  
    if(pid == 0) {  
        // child process code  
        sleep(60);  
    }  
  
    return 0;  
}
```

# References

- <https://shivammitra.com/operating%20system/zombie-and-orphan-process-in-opeating-system>
- Zombie and orphan process video - <https://youtu.be/L3YQDUuDjoo>
- <https://shivammitra.com/operating%20system/fork-exec-wait-in-operating-system/>